

Homework 7

Name: Put your name here

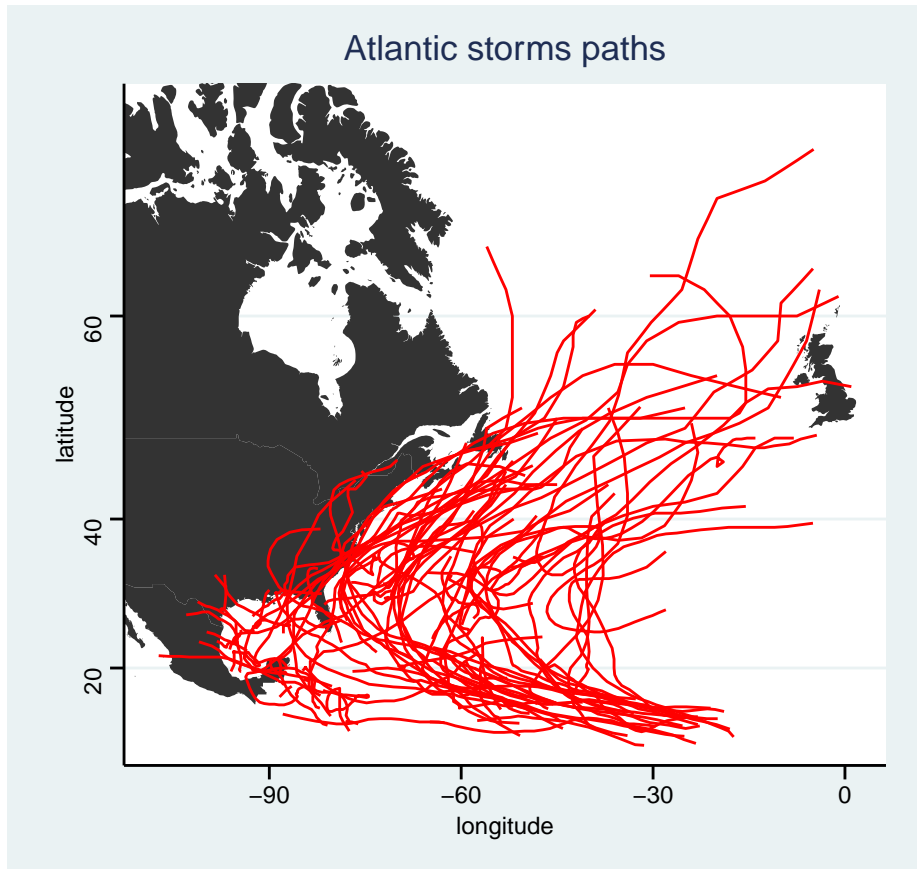
I worked with:

Click the “Knit” button in RStudio to knit this file to a pdf.

Problem 1: Storm paths

```
data(storms, package = "nasaweather")
ctry <- map_data("world",
                 region = c(
                   "usa",
                   "mexico",
                   "canada",
                   "uk"
                 ))
base_map <- ggplot(ctry) +
  geom_polygon(aes(x = long, y = lat, group = group)) +
  labs(
    x = "longitude",
    y = "latitude",
    title = "Atlantic storms paths"
  )

base_map +
  geom_path(data = storms, aes(x = long, y = lat, group = name), color = "red") +
  coord_map(xlim = c(min(storms$long), max(storms$long)),
            ylim = c(min(storms$lat), max(storms$lat)))
```



a.

answer:

b.

answer:

c.

answer:

Problem 2: Shiny App for Visualizing Baby Name Trends

In this assignment, as seen in class activity 19, you will extend a basic Shiny app to visualize trends in baby names. The foundation of the app includes a user interface for input and a server setup to filter data reactively. Your task is to expand this app to allow for the visualization of multiple baby names on the same plot, updating reactively with the press of an action button.

```
ui <- fluidPage(
  titlePanel("Baby Names Trend"),
  sidebarLayout(
    sidebarPanel(
      textInput("name", "Enter a Baby Name:", "Emma"),
      selectInput("gender", "Select Gender:", choices = c("Male" = "M", "Female" = "F")),
      actionButton("goButton", "Show Trend")
    ),
  ),
)
```

```

    mainPanel(plotOutput("nameTrend"))
  )
)

server <- function(input, output) {
  nameData <- eventReactive(input$goButton, {
    req(input$name) # Ensure the name input is not empty
    babynames %>%
      filter(name == isolate(input$name), sex == isolate(input$gender))
  })

  output$nameTrend <- renderPlot({
    req(nameData())
    ggplot(nameData(), aes(x = year, y = n)) +
      geom_line() +
      labs(title = paste("Trend for name", isolate(input$name)),
           x = "Year", y = "Number of Babies") +
      theme_minimal()
  })
}

shinyApp(ui, server)

```

Baby Names Trend

Enter a Baby Name:

Select Gender:

Male ▼

a.

answer:

b.

answer:

app url here

Problem 3: Spam using k-nn

This example looks at a data set of about 4600 emails that are classified as spam or not spam, along with over 50 variables measuring different characteristic of the email. Details about these variables are found

on the Spambase example on the machine learning data archive. The dataset linked to below is a slightly cleaned up version of this data. The only extra column in the data is `rgroup` which is a randomly assigned grouping variable (groups 0 through 99) which we will eliminate from the data.

Read the data in using the commands below to create a response `class` variable that contains the factor levels `spam` and `nonspam` with `spam` the first level.

```
# tsv = tab separated values!
spam <- read_delim("http://math.carleton.edu/kstclair/data/spamD.txt",
  delim="\t")

# some clean up
spam <- spam %>%
  mutate(class = fct_recode(
    spam,
    spam = "spam" ,
    nonspam = "non-spam"), # rename levels because caret doesn't like "non-spam"
    class = fct_relevel(class, "spam") # make "spam" the first level (our "positive")
  ) %>%
  select(-rgroup, -spam) # don't need random group variable and spam variable
levels(spam$class) # verify "spam" is level 1
## [1] "spam"      "nonspam"
```

a.

Answer: Your answer here

```
# Your code here
```

b.

```
set.seed(757302859) # set a seed
```

Answer: Your answer here

```
# Your code here
```

c.

Make a recipe for fitting k nearest-neighbor algorithm to the training data by inputting the formula and the preprocessing steps.

Answer: Your answer here

```
# Your code here
```

d.

Answer: Your answer here

```
# Your code here
```

e.

Answer: Your answer here

```
# Your code here
```

f.

Answer: Your answer here

```
# Your code here
```

g.

Answer: Your answer here

```
# Your code here
```

h.

Use the `tidymodels` package to do 10-fold cross validation as follows:

- use the 80% training data split from part b.
- tune your knn spam classifier based on accuracy
- consider neighborhood sizes ranging from size 1 to 31

Use the `results` to get the training set cross-validated estimates of the accuracy, precision, sensitivity and specificity of your final (“best”) classifier.

And use the following seed before running your `train` command:

```
set.seed(30498492)
```

Answer: Your answer here

```
# Your code here
```

i.

Answer: Your answer here

```
# Your code here
```

Problem 4: Incoming student characteristic

We will look at a “classic” college data set of a random sample of colleges and universities. To simplify our look at this data, we will filter to only look at MN, MA, and CA schools

```
colleges <- read_csv("http://math.carleton.edu/kstclair/data/Colleges.csv")
names(colleges)
## [1] "State"      "College"    "SATM"       "SATV"       "AppsReceive"
## [6] "AppsAccept" "HStop10"    "HStop25"    "FullTime"    "Tuition"
## [11] "RoomBoard"  "Books"      "Ratio"      "Donate"      "Expend"
## [16] "GradRate"   "Type"       "AvgSalary"  "NumFaculty"
colleges2 <- colleges %>%
  filter(State %in% c("MN", "MA", "CA"))
colleges2 %>% count(State)
## # A tibble: 3 x 2
##   State     n
##   <chr> <int>
## 1 CA      21
## 2 MA      19
## 3 MN      11
```

We will also just focus on student body characteristics (incoming class averages) for SAT and the HS variables (which are the proportion of the incoming class that is in the top 10% or 25% of their HS class). Here we select just these characteristics and college name and state.

```
colleges2 <- colleges2 %>% select(1,2,3,4,7,8)
colleges2
## # A tibble: 51 x 6
##   State College          SATM SATV HStop10 HStop25
##   <chr> <chr>          <dbl> <dbl>   <dbl>   <dbl>
## 1 CA    California Institute of Technolo 750  660    98    100
## 2 CA    California Lutheran University 495  436    23    52
## 3 CA    California Polytechnic-San Luis 547  455    47    73
## 4 CA    Chapman University          501  456    23    48
## 5 CA    Claremont McKenna College      670  600    71    93
## 6 CA    Harvey Mudd College            740  630    95   100
## 7 CA    Pitzer College                590  560    37    73
## 8 CA    Pomona College                700  640    80    98
## 9 CA    Scripps College               590  560    60    83
## 10 CA   Occidental College            570  510    52    81
## # i 41 more rows
```

Let's cluster schools by their incoming class characteristics.

(a)

Answer: Your answer here

```
# Your code here
```

(b)

Answer: Your answer here

```
# Your code here
```

(c)

Answer: Your answer here

```
# Your code here
```

(d)

Answer: Your answer here

```
# Your code here
```

(e)

Answer: Your answer here

```
# Your code here
```

(f)

Answer: Your answer here

```
# Your code here
```

(g)

```
library(GGally)    # install if needed
colleges2 %>%
  ggpairs(aes(color = cluster_km3),
          columns=c("SATM", "SATV", "HStop10", "HStop25"))
```

Answer: Your answer here

```
# Your code here
```